

DECISION THEORY FOR NON-CONSEQUENTIALISTS*

Wolfgang Schwarz

Draft, 2 November 2014

1 INTRODUCTION

Moral philosophers have said a lot about what acts would be right or wrong in decision problems for which all relevant facts are specified. However, any such account of “objective” moral status should arguably be accompanied by an account of “subjective” moral status: about what would be right or wrong relative to the limited information available to the agent.

For consequentialists, who measure the moral status of an act by the goodness of its consequences, a natural strategy is to adapt the framework of decision theory. If an agent doesn’t know which of her options would lead to the most good, she then ought to maximize the *expected* good – that is, the weighted average of all the good an act might bring about, weighted by the respective probabilities.

I will argue that the decision-theoretic approach to subjective moral status is not reserved for consequentialist theories. Properly understood, it does not presuppose that the rightness of acts is derivative of the goodness of states of affairs; it does not imply that moral agents should be risk-neutral; it can allow for agent-relative values and supererogatory options. Indeed, I will argue that in some respects the decision-theoretic approach is easier to justify in a non-consequentialist setting.

2 MORAL DECISION PROBLEMS

Consider a trolley problem. A runaway trolley is heading towards three people tied to the tracks. You can flip a switch that would redirect the trolley onto another track where it would run over a construction worker repairing the lights in a tunnel. What should you do? Different moral theories give different answers. Some say you should flip the switch because three deaths are worse than one. Others say you shouldn’t flip the switch, perhaps because killing one person is worse than letting three die.

* Ancestors of this paper were presented at the ANU and the University of Leeds. I’m indebted to the audiences on these occasions, as well as Christian Barry, Holly Lawford-Smith, and especially Seth Lazar for very helpful comments.

Now consider the following variation. A runaway trolley is heading towards three people tied to the tracks. You can flip a switch that would redirect the trolley to another track, where it *might* run over a construction worker scheduled to repair the lights in a tunnel. What should you do? This time, a crucial piece of information is missing. We don't know whether the construction worker is actually there and whether she would be harmed if you redirected the trolley. The missing information is especially important from the perspective of theories according to which you shouldn't flip the switch in the original problem. In the new problem, we don't know whether flipping the switch would amount to a killing or not. So even if we assume an uncompromising ban on killing, we can't say what you should do.

Is this a problem? Plausibly, moral theories don't need to give verdicts for arbitrary underspecified cases. "There are two buttons; which one should you press?" – That question can't be answered: we need more information about the buttons. Similarly, one might claim that in the second trolley problem, flipping the switch is obligatory if the construction worker is not in the tunnel and forbidden if the construction worker is in the tunnel, and that's all there is to be said. On this view, a moral theory only needs to provide an evaluation of fully specified decision problems, where all relevant facts are settled.

But this is an unsatisfactory position. In real life, we almost never have perfect information about all relevant facts. We don't know whether there are people on the tracks. We don't know how the economy will react to a stimulus plan. We don't know whether civilians would be harmed in a military operation. How should our actions be guided by our moral theory if our moral theory refuses to give a verdict until all factual uncertainty is resolved? And shouldn't endorsing a moral theory have some effect on one's choices – otherwise what's the point of morality? Parallel problems arise for our role not as agents but as advisers or judges. Suppose you ask me what you should do, perhaps because you know that I have more (but still incomplete) information. Shouldn't the moral theory I endorse manifest itself in my advice?

Discussions on this issue sometimes turn into debates about semantics. There is a sense of 'right' and 'ought' – often called *objective* – on which the right choice (what you ought to do) is the choice that is right in the light of all the facts, known and unknown. But there also seems to be a *subjective* sense on which an act is right if it is right relative to some contextually relevant information. For my topic, it isn't really important whether there is such a reading for deontic modals in English, or whether one of the readings is somehow prior to the other. What's important is that moral theories should not content themselves with verdicts about objective moral status.¹

¹ I will continue to use the entrenched labels "subjective" and "objective", but I should point out that my usage does not conform to every usage in the literature; in particular, my account of subjective moral status is not meant to satisfy the adequacy conditions Holly Smith puts forward in [Smith

I mentioned one motivation for caring about subjective moral status: that endorsing a moral theory should guide our actions and reactions. Another motivation (emphasised e.g. in [Jackson 1991]) is that there is something deeply troubling about agents who always make the objectively right choice without also happening to be omniscient. Such agents would constantly take incredible moral risks and always come out lucky. To illustrate, consider another variation of the trolley scenario, where the tunnel track is empty and you have a further option to stop the trolley by throwing a large rock in its way. Doing so would cause significant damage to the trolley. Objectively, it would then be best to flip the switch and let the trolley come to a natural stop on the empty track. But suppose you don't know that the track is empty. Suppose all you know is that one of the two tracks is empty and the other has three people tied to it. In that case, surely the only defensible choice is to throw the rock. A morally conscientious agent would not do the objectively right thing and flip the switch. Moreover, we can say what she would do without knowing to which track the three people are tied. In other words, we can evaluate the options in her decision problem without fixing all morally relevant facts. Our moral theory is not silent.

I should perhaps say a little more on what I mean by a 'moral theory'. A moral theory is a system of substantive moral principles, rules, guidelines etc. that determines a ranking of possible choices in decision problems. The ranking need not be complete, and it may be very coarse-grained, classifying some of the available options as right and others as wrong. The totality of our moral convictions is a moral theory. So is the totality of objective moral truths, if there is such a thing. Genuine moral disagreement is disagreement over moral theories.

I make no general assumptions about the internal structure of moral theories. In principle, a moral theory could be an endless, gerrymandered list of verdicts about specific decision problems, without any unifying principles. It is clear, however, that our own moral commitments do not take this form. At the other extreme, a moral theory might base all its verdicts on a single unifying principle – the categorical imperative perhaps, or the principle that pleasure is good and pain bad. A moral theory might also address decision problems rather indirectly, perhaps via character traits that ideal agents should possess. That alone would not yet constitute a moral theory in my sense since it would give no verdicts about decision problems. It turns into a moral theory once we specify how the relevant character traits manifest themselves in an agent's choices.

Let's say that the *domain* of a moral theory is the class of decision problems for which it provides a ranking of available acts; the domain together with the associated rankings is the theory's *extension*. A theory that deals exclusively with trolley problems would have a very limited domain. A theory that only deals with decision problems in which the outcomes of all options are fully specified also has a limited domain, although in a

2010], for reasons I will explain in section 8.

different respect. The above observations suggest that our moral commitments are more comprehensive than that. We can evaluate the options in decision problems without knowing exactly what would happen if a given option were chosen.

I will proceed on the assumption that credible moral theories should provide an evaluation of the options in a significant range of decision problems with imperfect information. I will also assume that the relevant kind of imperfect information can be modelled by a probability measure. If we're concerned with action guidance, the most salient interpretation of this measure has it stand for the agent's subjective degrees of belief. For example, if the agent is 80 percent confident that the construction worker is in the tunnel, the relevant decision problem will include a probability measure assigning probability 0.8 to that state of the world.²

It won't hurt if throughout this paper you interpret the probabilities in decision problems as the agent's degrees of belief. My own view is that moral theories should actually handle decision problems without fixing the interpretation of the probabilities. We can ask what an agent should do if she has such-and-such options and such-and-such information or evidence, without assuming that the agent, or anyone else, is aware of the information. At the limit, where we take into account all the facts, we get the "objective oughts" on which moral philosophers have traditionally focused. If instead we evaluate the options relative to the agent's degrees of belief, we get a kind of "subjective ought". If we evaluate them relative to the degrees of belief that would be appropriate given the agent's evidence, or relative to an adviser's or onlooker's degrees of belief, we get different kinds of "subjective ought".

That evidence or degrees of belief can be represented as a unique, coherent probability measure is of course an idealisation. More realistically, all we may have are unsharp qualitative constraints on probabilities: that A is much more probable than B , that C is approximately independent of D , etc. For the sake of simplicity, I will nevertheless focus on the idealized case in which the relevant probabilities are sharp. Cases with unsharp probabilities can then be treated by the method of supervaluation, but I will not discuss whether that treatment is fully adequate.

So the decision problems I will consider are given by a range of options and a probability measure over relevant possibilities. Moral theories should determine a ranking of these options. How should they achieve that? There is always the possibility of an endless, gerrymandered list. But suppose we want something more systematic. Can we use the framework of decision theory?

² In ordinary English, the term 'probability' arguably can't be used to mean degrees of belief. I use 'probability' in the mathematical sense in which any function that satisfies certain formal conditions is a probability measure.

3 THE DECISION-THEORETIC APPROACH

Here, in rough outline, is how standard decision theory evaluates an agent's options. Suppose the agent faces a choice between some options A_1, A_2, \dots . Suppose also that her degrees of belief are represented by a probability measure P over possible *states of the world* S_1, S_2, \dots . Each of these, together with any option, determines an *outcome*. Let $O[A_i, S_j]$ denote the outcome brought about by option A_i in state S_j . Suppose we have a utility function U (on which more soon) that assigns to every outcome a number. If we hold fixed a given option A_i , the utility function effectively maps each state S_j to a number $U(O[A_i, S_j])$: to the utility of choosing A_i in S_j . In the confusing jargon of probability theory, this mapping is called a 'random variable'. Its (equally ill-named) *expectation*, a.k.a. the *expected utility of A_i* , is the average of the values, weighted by their probabilities:

$$(1) \quad \sum_i P(S_i)U(O[A_i, S_i]).$$

Decision theory now ranks the available options by their expected utility in the sense of formula (1). The best options are those with greatest expected utility.

Decision theorists disagree on the interpretation of the utility function U . A popular approach in economics and psychology takes the function to measure the agent's personal wealth or welfare or well-being. Decision theory thereby turns into a theory of narrow-minded, self-interested *homines economici*. Unsurprisingly, the theory does not accurately predict real people's choices, and it is doubtful that it provides a plausible standard of rationality. From a Humean perspective, there is nothing irrational about agents who, say, care more about the well-being or ill-being of others than about themselves.

Whatever its relevance to a theory of rationality, the present approach lends itself naturally to consequentialist moral theories. We simply have to replace the personal utility function U by a *moral value function* V . Just as a rational *homo economicus* would promote her personal utility U , a moral agent should promote the moral value V . For traditional utilitarians, V could measure overall happiness in the world. For more contemporary consequentialists, it might be some kind of aggregate of welfare or preference satisfaction across all people and times. The advice to a morally conscientious agent would then be to maximize expected moral value. For the limit case where all information is given, this means to choose whatever option actually maximizes V . If only probabilistic information about the possible consequences of each option is available, the goodness of the possible outcomes (as captured by V) has to be weighted by the corresponding probabilities.

One problem with this approach concerns the value function V . Standard decision theory assumes that utility is a *cardinal* measure, meaning that it not only represents whether one outcome is better than another, but also by *how much* the first is better or

worse than the second. More precisely, it should be meaningful to say that an outcome O_1 is as much better than an alternative O_2 as a third outcome O_3 is better than O_4 . It is not obvious that every consequentialist moral theory can provide such a cardinal measure of goodness. To give an objective ranking of an agent's options for decision problems where all relevant information is given, a merely ordinal measure is enough.

Another problem is the justification of the expectation rule. Why should we interpret “promoting” moral value as maximizing its expectation? This is a substantive moral assumption. To be sure, there are long-run arguments in support of the expectation rule: if everyone follows the policy of maximizing expected moral value, then by the laws of large numbers it is almost certain that *actual* moral value will be maximized in the long run, compared to any other policy. But reflections on what would happen if everyone followed a given policy are not the kind of reason that should impress an *act* consequentialist. On the other hand, rule consequentialists will not agree that the right act is the one that maximizes expected moral value in the first place.³

One might try to justify the expectation rule by pointing out that it gives intuitively plausible verdicts in actual cases. But even that looks doubtful. For example, the expectation rule implies that agents should be risk-neutral about moral value: if an act A_1 would lead to 100 units of good, while A_2 would lead to either 0 or 200, each with probability 1/2, then the two options have equal expected moral value. But in real cases, they often don't seem equally good.

These problems only get worse if we consider non-consequentialist theories, which don't see the goal of morality in the promotion of some good attached to possible outcomes. To be sure, the discussion over “consequentializing” has shown that many non-consequentialist theories concerning objective moral status are extensionally equivalent to theories of a consequentialist structure. The trick is to individuate outcomes very inclusively so that the “outcome” of an act includes the act itself, as well as the conditions under which it was chosen. Any right- or wrong-making feature of acts can then be located in the corresponding outcomes, so that one can define a notion of goodness that is maximized if and only if an act is obligatory according to the original, non-consequentialist theory. But this does not vindicate the use of decision theory to tackle decision problems with imperfect information. For one thing, the cooked-up goodness measure tends to be merely ordinal, not cardinal.⁴ Moreover, risk-neutrality (and similar demands of decision theory) look even more dubious from a non-consequentialist perspective.

In response to these worries, one might try to adapt what is often regarded as the

³ The present point is raised by Alexander Pruss at <http://alexanderpruss.blogspot.de/2014/01/consequentialism-and-doing-what-is-very.html>.

⁴ By this I mean that any monotonic (order-preserving) transformation of the goodness measure would do just as well. Cardinal information is irrelevant if we only want to reproduce a theory's objective oughts.

orthodox account of subjective utility in modern decision theory. On this account, an agent’s utility function is not identified with any fixed measure of wealth or welfare. Rather, it is derived from the agent’s preferences over possible acts, as revealed by her choice dispositions. So-called *representation theorems* show that if the preference order satisfies certain qualitative conditions (known as decision-theoretic *axioms*), then it can be represented by a utility function U in such a way that the agent prefers an act A_1 to an act A_2 just in case the expected utility of A_1 is greater than that of A_2 .⁵

For the ethical case, this approach has the seeming advantage that we don’t need to postulate a prior goodness measure for outcomes, nor do we need to justify the expectation rule. All we need in order to vindicate the hypothesis that subjective rightness goes with maximizing expected value is the assumption that moral theories rank possible acts in a way that satisfies the axioms. The existence of a suitable moral value measure and the expectation rule then fall out from the representation theorems.

However, on closer inspection the preference-based approach also has a number of problems. First, it makes the application of decision theory in ethics utterly opaque. It would be nice if we could actually use the tools of decision theory to get clear about the right choices in, say, a trolley case with imperfect information. The preference-based approach takes such answers as input to determine a moral value function.

Relatedly, the present approach threatens to trivialise the idea that subjective moral status is determined in a decision-theoretic manner. The idea now reduces to the hypothesis that the moral ranking of acts in decision problems conforms to the decision-theoretic axioms. But arguably the axioms are too weak to characterize genuinely decision-theoretic moral theories.

In other respects, the axioms may be too strong – especially if we want to accommodate non-consequentialist theories. Recall that to make decision theory applicable to such theories, “outcomes” must be individuated very finely, ideally so that they include the acts that brought them about. It turns out that this move is incompatible with standard axiomatizations of decision theory in the tradition of von Neumann and Morgenstern [1944] or Savage [1954], for which it is essential that the very same outcome can be brought about by different acts.⁶

⁵ Classics in this tradition are [Ramsey 1931], [von Neumann and Morgenstern 1944] and [Savage 1954].

The preference-based approach has its roots in the positivist doctrine that scientifically legitimate quantities must be defined in terms of directly observable phenomena. As such, it is widely ridiculed in contemporary philosophy (e.g. [Meacham and Weisberg 2011]). However, it has a more respectable successor in the functionalist position, sometimes called “interpretivism”, defended e.g. in [Lewis 1974], [Stalnaker 1984], and [Blackburn 1998] (and indeed [Ramsey 1931]). That position takes an intermediary route between assuming subjective utilities as basic and deriving them from choice dispositions. I will advocate a similarly intermediary route for the case of moral value.

⁶ The assumption that outcomes are logically independent of acts is usually not presented as an axiom, but hidden in the definition of the act space. To see the importance of the assumption, consider the central *independence* axiom (also known as the *separability* axiom or the *sure-thing principle*) of

Finally, even if we grant that outcomes don't include the acts that brought them about, most of the "possible acts" over which preferences are assumed to be defined in classical axiomatizations do not correspond to options any agent could possibly face – let alone in a realistic, single decision situation. But all we can take for granted is that moral theories rank the options within a significant number of decision problems that agents may realistically face.⁷

In sum, neither of the traditional approaches to subjective utility looks promising if we want to apply decision theory to moral theories, including non-consequentialist theories.

There are further worries with the decision-theoretic approach, no matter how the moral values are defined. First, non-consequentialist theories often postulate agent-relative norms. Suppose some moral theory – call it M – says that in a certain situation Alice should dance with Bob, although Bob should not dance with Alice. How is this possible if the rightness of acts is derived from the impersonal goodness of outcomes? Either the outcome in which Alice dances with Bob is better than the one in which she doesn't dance with Bob, or it is not better. How could it be better relative to Alice and worse relative to Bob? One might say that different moral theories – different standards of goodness, different value functions – pertain to Alice and to Bob, but M is supposed to be a single moral theory, simultaneously giving verdicts for both Alice and Bob.

Second, the expectation rule seems to over-generate moral obligations. Suppose you think about going for a bike ride in the afternoon. If you go, there is a greater risk that you might get in an accident and injure (say) an innocent child than if you stay at home. Assuming that injuring the child has significant moral disvalue, staying at home plausibly comes out as having greater expected moral value. But are you really obligated not to go on the bike ride, merely due to the small risk that you might injure a child?

classical decision theory:

If two acts A and A' lead to the very same outcomes unless some condition C obtains, and if two further acts B and B' coincide in terms of outcomes with A and A' respectively whenever C obtains, and coincide with one another whenever C does not obtain, then the agent prefers A to A' iff she prefers B to B' .

If different acts can never lead to the same outcome, there are no acts at all that satisfy the stated conditions on A , A' , B , and B' : the axiom becomes empty.

One might try to accommodate non-consequentialist values (or, for the theory of rationality, non-consequentialist desires) without going all the way to include acts in outcomes, but it is an open question to what extent this is compatible with the spirit and/or the letter of classical axiomatic decision theories. There is a large literature on this topic; see e.g. [Tversky 1975], [Broome 1991: chs.5–6], [Pettit 1991], [Dreier 1996b], and [Joyce 1999].

⁷ Arguably the best response to this problem (advocated e.g. in [Joyce 1999] for rational choice and in [Broome 1991] for consequentialist moral theories) is to interpret the preference relation not directly in terms of choice behaviour, but as comparative desirability, which is only partly and indirectly related to choice. But it is not obvious that every credible moral theory must provide a sufficiently rich relation of comparative moral desirability.

Relatedly, if the subjective moral imperative is to maximize expected moral value, there seems to be no room for supererogatory options: options that are morally better than the alternatives and yet are not morally required. Similarly, there seems to be no room for moral dilemmas in which all options would be wrong. Non-consequentialists often want to allow for supererogation and moral dilemmas.

Finally, it has often been pointed out that as a practical guide, the principle to maximize expected moral value is pretty useless, as it seems to presuppose not only that the decision maker is conscious of the relevant probabilities and moral values, but also that she is capable of performing the required calculations, which quickly becomes intractable if the space of relevant options and states is as large as it usually is.

I will argue that all these worries can be answered. As a first step towards an answer, I suggest that we replace the classical framework of “expected utility theory” with something a version of Richard Jeffrey’s [1965] “Logic of Decision”.

4 THE LOGIC OF DECISION

The Logic of Decision [Jeffrey 1965] introduced some key innovations to decision theory. One is to jettison the traditional distinction between acts, states, and outcomes. In Jeffrey’s account, probabilities and utilities (or “desirabilities”) are both defined over entities of the same kind: propositions. Among other things, this has the advantage that one can consider probabilities conditional on acts, or probabilities of complex propositions involving acts – something that isn’t possible in earlier frameworks like [Savage 1954] and that proves important to give a clear treatment, for example, of Newcomb problems (see [Joyce 1999: 117]). Moreover the basic utilities in Jeffrey’s account are assigned not to restricted “outcomes”, but to maximally specific propositions, which allows us to easily capture agents who care not only about outcomes but also about how the outcomes were caused.

Jeffrey’s theory is also beautifully simple. Let \mathcal{A} be the set of propositions whose truth-value is relevant to a given decision problem. Assume that \mathcal{A} is closed under conjunction, disjunction and negation. Every element of \mathcal{A} is then entailed by a *maximal* element: a proposition that entails A or $\neg A$ for every A in \mathcal{A} . Let \mathcal{W} be the set of these maximal elements. Intuitively, a proposition in \mathcal{W} is a description of a possible way things could be that includes everything that matters for the evaluation of the options. Assume we have a value function V over the members of \mathcal{W} , and a probability measure P over \mathcal{A} . Then we can define the *desirability* of any proposition $A \in \mathcal{A}$ as the (conditional) expected value

$$(*) \quad EV(A) = \sum_{W \in \mathcal{W}} P(W/A)V(W).$$

Here $P(W/A)$ is the probability of W conditional on A , which equals the ratio $P(W\&A)/P(A)$ whenever $P(A)$ is greater than 0.⁸

Here is a simple example. Suppose you wonder whether to have a glass of wine with dinner. The problem is that wine sometimes gives you a headache, and you don't like having headaches. Suppose there is nothing else you care about to which your choice might make a difference. We can then take the atomic propositions in \mathcal{A} to be the proposition that you have a glass of wine (for short, W) and the proposition that you will get a headache (H). The set \mathcal{W} of maximally specific propositions then has four elements: (1) $W\&H$, (2) $W\&\neg H$, (3) $\neg W\&H$, and (4) $\neg W\&\neg H$. Conditional on the hypothesis that you have wine, (3) and (4) have probability zero: given that you do have wine, it is certainly not the case that you don't have wine. So the desirability of having wine is given by

$$Des(W) = P(W\&H/W)V(W\&H) + P(W\&\neg H/W)V(W\&\neg H),$$

which can be further simplified to

$$Des(W) = P(H/W)V(W\&H) + P(\neg H/W)V(W\&\neg H).$$

That is, to evaluate the desirability of having wine, we have to consider two scenarios: you might have wine and get a headache ($W\&H$), or you might have wine and not get a headache ($W\&\neg H$). The desirability of having wine is the average of the values assigned to these scenarios, weighted not by their absolute probability, but by their probability conditional on a decision to have wine. The desirability of not having wine is similarly given by

$$Des(\neg W) = P(H/\neg W)V(\neg W\&H) + P(\neg H/\neg W)V(\neg W\&\neg H).$$

Jeffrey suggests that desirability is a measure of choiceworthiness: rational agents should choose options with greatest desirability. This is a natural assumption. However, it is widely (though not universally) taken to go wrong in unusual scenarios like Newcomb's problem where an act would be evidence for a desirable or undesirable proposition without having any causal bearing on that proposition (see e.g. [Gibbard and Harper 1978], [Joyce 1999: sec. 5.1]). According to proponents of "causal decision theory", an adequate standard of choiceworthiness should somehow take into account the causal relationship between a possible choice and the relevant scenarios in \mathcal{W} . Different proposals have been made to render this idea precise – classical treatments include [Gibbard and Harper

⁸ The present definition of desirability is unproblematic if the relevant set \mathcal{W} is finite. Standard ways of extending the notion of a sum to infinite sets are often employed to deal with decision problems that require larger sets of possibilities, but it is not entirely unproblematic that these extensions are adequate, and they don't work for all cases. Fortunately, most real-life decision problems can be modelled with a finite set \mathcal{W} .

1978], [Lewis 1981], [Skyrms 1984] and [Joyce 1999]. I will not enter into these details here. In everyday decision problems, ranking options by their desirability (in Jeffrey's sense) almost always leads to sensible results.

Jeffrey's approach has not been popular among economists and psychologists, perhaps because it makes no testable predictions at all about what rational agents will do, independently of their desires. This is a consequence of the fact that the ultimate bearers of value in Jeffrey's system are the proposition \mathcal{W} that specify not only an act's consequences, but also the very act itself, as well as the history that precedes it. Hence Jeffrey's theory does not predict that if you would choose an apple over a banana, and a banana over a carrot, then you would choose the apple over the carrot. You might well have a desire for apple-if-the-alternative-is-banana, but no desire for apple-if-the-alternative-is-carrot. Similarly, it does not say that you should prefer a gamble between \$10 and \$20 to a sure outcome of \$1: if you hate gambles, you may well prefer the sure dollar. For analogous reasons, Jeffrey's theory makes no predictions about the scenarios of Allais [1953] and Ellsberg [1961] that seem to provide counterexamples to traditional expected utility theory. Indeed, any pattern of choice dispositions whatsoever can be made to trivially conform with the standard of maximizing desirability: if the agent is disposed to choose option O in a decision problem D , simply assume that she assigns high basic value to elements of \mathcal{W} that entail that the agent chooses O in D .

There is an elegant representation theorem for Jeffrey's theory, due to Bolker [1966], which shows how the value function V can be derived from a qualitative preference relation,⁹ but the relevant preference relation is not a relation between possible acts; rather, it is a relation of comparative desirability between arbitrary propositions in \mathcal{A} . The Bolker-Jeffrey axioms do not constrain an agent's possible choices.¹⁰

In my view, all this is just as it should be if we're interested in a pure theory of *instrumental* rationality: in a theory that says what an agent should do to further her goals, without presupposing anything about these goals. Substantive predictions about what such an agent should do must always be based on assumptions about the agent's goals.

⁹ Like Ramsey [1931] and Savage [1954], Jeffrey and Bolker actually consider the harder task of deriving both a utility function and a probability measure from the agent's preference relation. For our purpose, we can take the probabilities as given, which simplifies the problem.

¹⁰ To illustrate, consider the following axiom of *averaging*:

If two propositions X and Y are incompatible and X is preferred to Y , then X is preferred to the disjunction $X \vee Y$, which in turn is preferred to Y .

The idea is that an unspecific proposition $X \vee Y$ cannot be better or worse than both of the more specific properties X and Y . Killing an innocent cannot be worse than killing an innocent man and also worse than killing an innocent woman. This is compatible with the assumption that killing an innocent whose gender is unknown is worse than killing an innocent known to be male and also worse than killing an innocent known to be female.

In any case, I want to suggest that a theory along Jeffrey’s lines is an adequate framework to address moral decision problems. To that end, it will be useful to adopt a minor variation of Jeffrey’s framework, due to Lewis ([1979], [1981]), on which probabilities and utilities are defined not over propositions but over *properties* – things like living in London, liking lasagna, or being friends with a plumber.

Assigning probabilities to properties may seem odd. What does it mean to say that there is a 50 percent probability of living in London? Well, recall that the probabilities we are interested in capture imperfect information about a decision problem. In this context, to say that some property F has probability x simply means that the agent in question has property F with probability x . Using properties as bearers of probability is useful to model cases in which an agent is uncertain who they are or what time it is (see [Lewis 1979]). For our topic, the more important advantage of properties is that they provide a natural locus of agent-relative and time-relative value (compare [Lewis 1989: 73–76], [Dreier 1996a]).

To see how, return to the above case where (according to some moral theory) Alice should dance with Bob while Bob shouldn’t dance with Alice. This is hard to explain if the bearer of moral value are propositions or states of affairs. Is it good or bad, right or wrong, if Alice dances with Bob? It looks like the situation is right *for Alice* and wrong *for Bob*. That is easy to capture if the bearers of value are properties: in the given situation *dancing with Bob* is right for Alice, while *dancing with Alice* is wrong for Bob. There is only one state of affairs, but it involves different properties for the two agents.

Having properties as bearers of value also allows us to apply decision theory without first consequentializing moral theories. After all, act types are properties. So we can directly assign moral value to things like intentionally killing an innocent, or resisting a strong desire to gossip, without any detour through finely individuated outcomes.¹¹

If the algebra \mathcal{A} over which probability and desirability are defined is an algebra of properties, then its maximal elements \mathcal{W} are maximally specific properties: properties that entail, for every other property in \mathcal{A} , whether it is instantiated or not. I will call such maximal properties *profiles*. Intuitively, the profile brought about by a choice in a concrete decision situation is the totality of all properties that come to be instantiated as a result of the choice. If we can find a moral value function V that provides a cardinal ranking of possible profiles, we can use Jeffrey’s formula (*) to evaluate the options in decision problems with imperfect information.

¹¹ An alternative response to the worry that basic value in non-consequentialist theories does not attach to states of affairs would follow the observations of [Pettit 1991] and argue that the utilities taken as basic in standard decision theory are actually derived from more basic utilities assigned to properties instantiated in the relevant outcomes. But that alone would not allow for agent-relative and time-relative values.

5 MORAL VALUE

I mentioned above that for any pattern of choices there is a trivial way to construct a value function V that makes the choices conform to the principle of maximizing expected value in the sense of (*): we simply assign high value to making the relevant choice in the relevant decision problem. However, I want to defend a more substantive, non-trivial role of decision theory in moral theories. The basic idea is that we can use decision theory to systematically turn the verdicts a moral theory gives for decision problems with perfect information into verdicts about decision problems with imperfect information – to derive a theory’s “subjective oughts” from its “objective oughts”.

A theory of objective moral status effectively provides a (partial) ranking of what I called profiles. For example, it might say that flipping the switch in such-and-such a fully specified trolley problem is better than not flipping the switch. If the ranking is sufficiently complete and sufficiently cardinal (as argued below), it can be represented by a value function V that allows us to rank the options in ordinary decision problems by their expectation.

I should emphasize that the ranking of profiles, and the value function that represents it, is not supposed to be a moral primitive. I do not assume that promoting V is the goal of moral agency. Instead, a moral theory might attribute basic moral relevance to specific properties such as keeping promises, harming innocents, obeying God’s commands, or acting in accordance with the categorical imperative. In a given situation, several morally relevant properties may be instantiated together, in which case a moral theory might say how they combine to determine an overall ranking of a profile. For example, a theory might say that under certain conditions one is in fact required to harm an innocent, if that is the only means to prevent a tragedy. Even if a moral theory evaluates acts by their relation to some independently given measure of good, that measure may or may not be represented by the value function V . For example, a rule utilitarian might take overall happiness as the ultimate standard for evaluating acts, without ranking profiles by the amount of overall happiness they entail. (Rather, it will base its ranking on whether a profile involves an act type whose general adoption would lead to greater happiness in the world.)

Any moral theory that provides an objective ranking of options in fully specified decision problems thereby provides a partial ranking of profiles. What we do not have at this stage is a ranking of arbitrary propositions, or of possible acts across arbitrary decision problems. So we cannot use the representation theorems of axiomatic decision theories to argue that if that ranking satisfies such-and-such conditions, then it is represented by a cardinal value function etc. We have more work to do.

It is common in decision theory to assume that the objective ranking of profiles (or outcomes, in Savage-type theories) is *complete* in the sense that it says of any two profiles

W_1 and W_2 either that W_1 is better than W_2 , or that W_2 is better than W_1 , or that both are equally good. It is not plausible that every moral theory should provide a complete ranking in this sense. Incompleteness can arise in several ways. First of all, moral theories may have little or nothing to say about the comparative status of profiles that do not involve a single antecedent decision situation. Second, moral theories are often vague and indeterminate in various respects. For example, if we say that benevolence is good, without giving a perfectly precise definition of benevolence, our theory may fail to rank profiles involving borderline cases of benevolence. Third, a moral theory might define different dimensions of right and wrong, but no general rule for aggregating these into a single moral ranking. Fourth, a moral theory may explicitly *allow for* a whole range of rankings. For example, a moral theory might say that if the only way to prevent a tragedy is to harm an innocent, no choice is straightforwardly right, nor are the two options exactly equal in terms of rightness; a morally conscientious agent could go either way, depending on how she personally resolves that difficult choice.¹²

Incompleteness of the first type is irrelevant for our application, which only requires well-defined values for the profiles within any given decision problem. The other forms of incompleteness can all be modelled by assuming a whole class of moral value functions. If an option A maximizes moral value relative to some legitimate value functions and not others, we can then say that there is no fact of the matter about whether you ought to do A . Alternatively, if the multitude of value functions represents indeterminacy of the fourth kind, we can say that A is not required but permitted. As I will further explore in section 7, having a class of value functions nicely captures the intuition that morality often leaves open exactly what we should do.

Standard decision theory also requires the basic value functions V to be cardinal, as otherwise the arbitrary choice of numerical representation will affect the ranking of options in decision problems with imperfect information. To some extent, we can actually allow for such representation dependence. It will generate even more indeterminacy of subjective moral status. But of course we do need at least *approximately cardinal* structure if we want to give non-trivial verdicts about decision problems with imperfect information.

Fortunately, it is plausible that our moral evaluations of profiles has approximately cardinal structure. It is hard to see how we could even start addressing decision problems with imperfect information if we didn't have such cardinal evaluations. Take an example. You are on a boat and notice what seems to be a white shark approaching a group of surfers. You have to decide whether to kill the animal with a harpoon. Let's say you are 80 percent certain that the fish is a dangerous white shark, and give 20 percent credence to the possibility that it is a harmless, but also endangered species of shark. What should

¹² The non-obvious fourth possibility was brought to my attention by [Gert 2004] and [Rabinowicz 2008], where it is actually used to model a form of the third possibility, what Chang [2002] called *parity*.

you do? Here are the most relevant things that might happen, depending on your choice:

- (1) You kill a white shark that would otherwise have attacked the group of surfers.
- (2) You kill a harmless and endangered fish.
- (3) You let a white shark attack the group of surfers.
- (4) You let a harmless fish approach the group of surfers.

To simplify, let's say that shooting would come with an 80 percent probability of (1) and a 20 percent probability of (2). Not shooting would come with an 80 percent probability of (3) and a 20 percent probability of (4). What makes your decision problem non-trivial is that (1) and (4) are much better than (2) and (3). It also matters that (3) is worse than (2), although both are deplorable. Moreover, it matters that (3) is *a lot* worse than (2), whereas (1) is not a lot worse than (4), by comparison with (3) and (2). If our moral theory doesn't allow us to make these quasi-cardinal comparisons, it is unclear how it could give sensible advice about what you should do. But if we have verdicts like these, we can meaningfully apply decision theory.

There are also independent reasons to think that our moral theories should provide fine-grained, approximately cardinal evaluations of profiles. I mention three.

First, the features of acts that matter to morality are almost never clear-cut, and often come in degrees. Imagine a sorites sequence leading from a clear case of murder to a clear case of innocent behaviour. (For example, gradually change the victim from an ordinary human to a robot, or a human in a final vegetative state, or a loose assembly of amino acids.) It is implausible that at some point, a minute difference between two cases makes an enormous moral difference – that one act in the sequence is simply forbidden and the next one permissible. Or consider all the morally relevant acts that come in degrees. If helping someone is right and harming her wrong, then surely helping her more is better than helping a little, and harming more worse than harming a little.

Second, moral theories should somehow be connected to our practice of blame and praise, punishments and rewards, and these come in “cardinal” units. We tend to criticise people less for cheating on their taxes than for sexual assault. It is implausible that this does not track any moral difference between the acts.

Third, we often need to balance moral considerations against other moral and non-moral considerations. Most of us think that extra-ordinary circumstances can make it permissible to lie or even kill an innocent if that is the only way to prevent a tragedy. Moreover, we think that people's obligations are sensitive to the amount of sacrifice fulfilling the obligations would require. Some obligations justify more sacrifice than others. Either way of balancing would make little sense if we didn't have approximately cardinal moral values.

Finally, if our objective moral theory contains insufficient cardinal information, we can take a lesson from the preference-based approach and work backwards: consider the

verdicts we want for decision problems with imperfect information and try to retrofit the cardinal structure of objective values. To illustrate, imagine some scenario in which an agent faces a choice of letting one person die or taking a 50 percent chance that either two people or no-one will die. At first glance, it might seem that all we can say from an objective point of view is that in the given scenario letting two people die would be worse than letting one die, which again would be worse than letting no-one die. This merely ordinal ranking is not enough to determine what the agent should do, if she doesn't know what would happen as a result of taking the risk. But suppose we also think that it would be wrong for the agent to take the risk. Assuming that acts are evaluated in terms of expected moral value, it follows that the difference in moral value between the profile in which two people die and the profile in which one dies is less than the difference between the latter and the profile in which no-one dies. Can we make sense of these cardinal judgements about profiles? Arguably yes. For example, the judgements could be explained by the assumption that there is something intrinsically wrong not only with letting people die but also with imposing risks on people. This wrong is involved in the profiles with zero and two deaths, but not in the profile with one death, which explains why the latter gets ranked higher than it would if the number of deaths were all that mattered.¹³

There is no guarantee that the forwards and backwards route will always cohere. Consider an “absolutist” moral theory according to which it is never permissible to (say) kill an innocent, no matter the circumstances. This is a claim about objective ought, since it doesn't say what to do in situations where it is uncertain whether a given act would involve a killing of an innocent. Absolutists might suggest that one should never act in a way that might possibly involve a killing of an innocent. However, with some imagination it is not difficult to see that almost every option in almost every real-life decision problem would thereby come out as impermissible. So consider a moderate absolutist who agrees that it is sometimes permissible to incur very low probabilities of killing an innocent. So far, this combination of objective and subjective claims is easy to model in the decision-theoretic framework. But suppose the absolutist's account of objective moral status includes more detailed, cardinal verdicts: it says that killing someone is always worse than letting people die, and that the moral difference between letting $n + 1$ people die and letting n people die is the same for all n . Now we have a problem.

In fact, we have several problems. A first is how to model these verdicts about objective status by a numerical value function at all, even without turning to cases with imperfect

¹³ Personally, I think it is an open question whether the agent's degrees of belief should factor in a profile's moral value. We certainly criticise people for acting in a way that might have caused serious harm, even if no actual harm occurred, but those judgements might well be judgements about subjective moral status. See [Smith 2010: 87–89].

information. If killing an innocent has finite disvalue and there is a fixed amount x by which letting $n + 1$ people die is worse than letting n people die, then it looks like for some n , letting n people die will have to be worse than killing an innocent. A popular way to prevent this conclusion is to say that killing an innocent should have infinite disvalue. Alternatively, we could move to a non-Archimedean number system in which there are finite limits of $n \cdot x$ as n goes to infinity (see [Bernstein and Wattenberg 1969]); or we could use vectors of reals instead of single numbers to reflect the apparently “lexical” ordering of profiles. There is certainly no conclusive reason why legitimate value functions should only take real-numbered values. The basic decision-theoretic approach does not presuppose that they do.

Having dealt with this preliminary problem in some such fashion, we can now see the tension between the moderate absolutist’s claims about objective and subjective status. If killing an innocent is “infinitely worse” than letting someone die, it looks like, by decision-theoretic rules, it is never permissible to, say, rescue a life if that comes with a positive probability of killing an innocent; and then it is hard to see what other acts could be permissible. Thus, if we work backwards from the verdicts our moderate absolutist wants to give for decision problems with imperfect information, it looks like we get objective moral rankings that contradict the absolutist’s direct cardinal judgements about profiles.

There are many ways in which one might try to resolve this apparent decision-theoretic incoherence, apart from reconsidering the relevant judgements. Two obvious strategies are to identify further relevant features of profiles (following the above treatment of risk) or to revise the implicit assumption that the overall value of a profile is simply the sum of the value of its components. These possibilities would be worth exploring before we conclude that the moderate absolutist position is incompatible with the decision-theoretic approach.¹⁴

Before we continue, it may be worth emphasizing that the use of numbers to represent moral value is a matter of mathematical convenience. By suggesting that moral theories should use the framework of decision theory, I don’t suggest that they should include a specific numerical value assignment to possible profiles (or a class of assignments). We don’t need to decide whether flipping the switch in a trolley problem has moral value 12.7 or -4.5. The proposal is rather that (a) moral theories should provide a ranking of profiles that is approximately cardinal in the sense that one can say whether W_1 is a lot, or only a little, better than W_2 , and that (b) the evaluation of decision problems with imperfect information should be systematically determined by the ranking of the profiles, in the sense that an option is ranked as better than another iff it has greater expected moral value relative to every value function V that respects the cardinal ranking of profiles.

¹⁴ Some attempts at modelling absolutist theories in a decision-theoretic setting are discussed in [Jackson and Smith 2006] and [Colyvan et al. 2010].

The actual numbers cancel; they need not figure anywhere in the moral theory.

6 EXPECTATIONS

Since Jeffrey-type decision theories put no substantive constraints on patterns of choice, any (ordinal) ranking of options in decision problems is a ranking by expected moral value in the sense of (*) relative to some moral value function. The decision-theoretic approach becomes non-trivial only if we impose constraints on the kinds of properties that enter into the determination of moral value. In particular, the more the value of a profile depends on features that may be unknown to the decision maker, the less trivial the evaluation by expected utilities.

The expectation rule becomes highly non-trivial if the moral value function is identified with some prior measure of good, as it often is in consequentialist theories. As I mentioned earlier, it is far from obvious that the right way to promote a good in the light of limited information is to maximize its expectation. Perhaps the most promising line of response to this challenge is the one pursued by Broome [1991]. Adapted to Jeffrey-type decision theories, Broome's basic idea is that once a probability measure is given, the consequentialist goodness ranking pertains not only to complete profiles, but to arbitrary propositions. If the ranking then satisfies the Jeffrey/Bolker axioms, it can be represented by a moral value function relative to which the ranking of propositions goes by expected value in the sense of (*). And so we get, among other things, a ranking of the propositions that represent the agent's options, and it will go by expected value in the sense of (*).

I will not discuss the prospects for this strategy, since I am more interested in the decision-theoretic approach for non-consequentialist theories. Here the situation is easier. As explained in the previous section, if moral value is not identified with some external, prior good, the construction of the value function can be sensitive to verdicts we want the theory to give for decision problems with imperfect information. In contrast to decision-theoretic forms of utilitarianism, we can, if we want, punish risk or reward fairness. The substantive assumption is no longer the expectation standard itself, but the assumption that objective value often turns on features of which the decision maker is unaware, and that objective and subjective value judgements should cohere in the sense of the previous section. These are the assumptions that make the decision-theoretic approach non-trivial.

I have no knock-down argument that any credible moral theory should satisfy these assumptions. But it does seem plausible that the moral evaluation of our options often turns on different possibilities about which we are not fully informed. This is what makes decision-making with imperfect information hard. So we need some way of taking into account the various possibilities. The expectation rule provides a simple, unified and

systematic way to do that. It can turn any value function for profiles into rankings for almost every decision problem an agent could face.

This systematicity is a great advantage. The space of possible decision problems is enormous. There are countless ways to be uncertain about the number of people on the tracks in a trolley problem; there are many more possibilities if we also take into account who these people might be; you may be further uncertain whether you face a trolley problem in the first place; you may be uncertain whether the switch is broken, whether the information you were given is correct, and so on. None of these cases is absurdly far-fetched or unrealistic. If we want our moral theory to give verdicts for a substantive range of decision problems people actually encounter, and if we don't want it to do that by giving an infinite list of *ad hoc* decision rules for individual problems, we need a unified set of principles.

To be sure, there are systematic alternatives to the expectation rule. But to the extent that they diverge from the expectation rule, these alternatives tend to be extremely implausible. Consider the *maximin* rule, which says that one should choose an option with the best worst-case profile. If a medical treatment has a 99.9 percent chance of curing a crippling condition, but a 0.1 percent chance of making it slightly worse, maximin suggests that it would be wrong to prescribe the treatment (because the worst-case outcome, the slightly worsened condition, is worse than the outcome of non-action). Or consider the popular, but unnamed strategy to do whatever is most likely to be best. This suggests that you should prescribe a treatment that has a 51 percent chance of slightly improving a minor condition and a 49 percent chance of killing the patient.

Another class of alternatives that are sometimes put forward by non-consequentialists are *threshold rules* according to which an act is right or wrong if it is sufficiently probable that certain states of affairs obtain – say, that you should prescribe a treatment only if it is more than 95 percent likely to cure the patient. Such rules involve hard to justify cut-offs, they don't generalize well, and tend to have implausible consequences for various edge cases. For example, the rule just outlined does not extend easily to cases where your choice is to prescribe the treatment to several people at once, and it wrongly assumes that it is irrelevant what happens in the 5 percent cases where the cure doesn't work.¹⁵ An adequate rule should consider all things that might happen, and it should take them into account in proportion to their probability. This is just what the expectation rule does.

Here is another reason to think that the expectation rule should play a role in the moral evaluation of decisions. Above I mentioned that a Jeffrey-type decision theory is plausible as a theory of purely instrumental rationality. It says how to rationally pursue one's desires or goals in the face of limited information, without assuming anything about

¹⁵ These and other problems for threshold rules have been extensively discussed elsewhere, see e.g. [Feldman 1995], [Jackson and Smith 2006] and [Lazar Unpublished].

the content of these desires or goals. It is natural to think of objective moral judgements as putting forward moral goals – not to kill innocents, etc. Thus the corresponding application of decision theory says how one should rationally pursue these moral goals in the face of limited information. Arguably, this is exactly what we are looking for when we want to extend our moral theory to cases of imperfect information.

Relatedly, imagine an agent who cares about nothing but doing what’s morally right. Arguably, such an agent could still face non-trivial decision problems. And it would be odd if she would then have to make choices that violate the principles of a minimal, Jeffrey-style decision theory. Morality does not require practical irrationality.

As I said, I have no knock-down argument that any credible moral theory should have a non-trivial decision-theoretic structure of the kind I have advocated. Consider the view that the only right- or wrong-making features of acts are located in the agent’s intentions, and suppose people generally know what intentions would go along with the choice of a given option (which is plausible if the options actually *are* intentions, as argued e.g. in [Weirich 1983] and [Sobel 1983]). Then there is (at least on the surface, and from the decision-maker’s own perspective) never any problem about decision-making with imperfect information, since agents are always fully informed about the morally relevant features of their options. For another example, consider the view that the moral status of options is given by their evidential expected value as given by (*), relative to some cardinal objective value rankings, while the rational status of options is given by causal expected value as defined by some form of causal decision theory. There is obvious decision-theoretic structure in such a theory, but it is not quite the structure that I think is most plausible for moral theories.

7 OPTIONALITY AND SUPEREROGATION

In these final two sections I want to address the last two worries left from section 3: the problems of optionality and supererogation, and the problem of practical guidance.

The problem of supererogation is to explain how it can be morally permissible to choose an option even though another option would be morally better. It is hard to see how this could be true if (a) we are always morally obligated to choose an option with greatest expected moral value, and (b) an option is morally better iff it has greater expected moral value.

More generally, the decision-theoretic approach seems to leave us little freedom to follow our personal goals and whims as long as these don’t have serious moral repercussions. If morality demands that we choose options with greatest expected moral value, the only choices left to us would concern cases where several options are exactly tied for expected value (or cases where the expectations are undefined). It is implausible that such cases

are very common, and they certainly don't sustain the kind of personal freedom many of us think we have.

However, remember from section 5 that the decision-theoretic approach doesn't assume a unique moral value function (or a value function that is unique up to positive linear transformations). Moral theories don't need to provide a complete ordering of all possible profiles. They may not fix how exactly different right-making and wrong-making features should be balanced; they may even explicitly delineate a range of permissible rankings. As a result, we have to evaluate an agent's options by a whole range of value functions. While ties relative to any single value function will be unlikely and rare, it is not unrealistic that different permissible value functions will often rank an agent's options differently. If we say that an option is obligatory only if it maximizes expected moral value relative to all permissible value functions, we get a considerable amount of personal freedom.

We still don't get supererogation. The central feature of supererogatory acts is that they are determinately better, yet not obligatory. A superficial way to allow for that is to dissociate obligatoriness from maximizing expected value. Thus we could follow the idea of "satisficing consequentialism" and say that one is obligated to choose an option only if it is (determinately) *a lot* better in terms of expected moral value than the alternatives, or if it is the only option whose expected value (determinately) exceeds a certain threshold, relative to some scale. This is still an essentially decision-theoretic account insofar as the moral status of an agent's options is still determined by their expected moral value.

Comparing an option's expected value to fixed points on the value scale might also provide a means to allow for moral dilemmas in which all options would be wrong. They would be wrong in the sense that their expected moral value lies below the fixed threshold that makes an option wrong.¹⁶

I have no general objections to these proposals, but I don't think they suffice to fully account for the phenomenon of supererogation. Consider an everyday scenario in which you could use a significant portion of your yearly income to rescue a friend from losing their home, or to save many distant strangers from starvation, and suppose we judge this act to be supererogatory: laudable, but not obligatory. Crucial for this judgement is clearly that rescuing the friend or the strangers would involve a great personal sacrifice. If you didn't care about your income and had no other use for it, helping would be obligatory. Thus in cases where you're allowed to keep the money, your personal reasons for keeping it somehow trump the moral reasons for giving it away. Moreover, the trumping is a *moral* kind of trumping: it's not that you are morally required to give away the money, but all-things-considered allowed to keep it. Whether or not you're allowed

¹⁶ Another type of moral dilemma in decision-theoretic accounts are so-called unstable decision problems, where any given choice makes a different choice preferable. Some forms of "ratificationism" entail that no act is permissible in these situations (see e.g. [Harper 1984]).

to keep the money is a genuinely moral question, something different moral theories will disagree about. On the other hand, it would be wrong to count your reasons for keeping the money as trumping *moral* reasons, for then giving away the money would be morally wrong.

Genuine supererogation therefore requires that non-moral reasons can morally override moral reasons (see [Portmore 2008]). This may seem puzzling, and some might want to conclude that the idea of supererogatory acts should simply be given up. However, when we turn to decision problems with imperfect information, the present phenomenon does seem to be fairly common. How can it be morally acceptable for me to go on a bike ride, given that there's a tiny chance that I will thereby cause harm to innocents?¹⁷ It looks like the non-moral goods probably involved in going on a bike ride can override the low risk of harming an innocent. Here again, the crucial point is that going on the bike ride isn't just "all things considered" permissible; it is *morally* permissible.

[?] develops a promising approach to model genuine supererogation. His focus is on objective moral status and consequentialist theories, but the basic strategy also works in our more general context.

To adequately account for supererogation, a moral theory should recognize nonmoral values or reasons, and specify how these may be balanced against moral reasons. In essence, we need not one but two rankings of profiles. The first ranks profiles by their purely moral status, the second ranks them in terms of the agent's well-being, welfare, preference satisfaction, or whatever non-moral considerations are acknowledged by the theory. (What kinds of non-moral reasons can trump moral reasons is again a moral question that we can't fix once and for all, for every moral theory.) Finally, the theory must say how the two kinds of value may be balanced against each other to determine an overall ranking. A simple way to model this might define a range of permissible overall value functions as weighted averages of the moral and non-moral value functions, relative to some fixed scale. The agent's options can then be evaluated in terms of expected overall value relative to all permissible overall value functions.¹⁸ An option is morally best iff it has greatest expected moral value. But the agent is not always obligated to choose the morally best option because the option may no longer be best when non-moral values are taken into account. In that case, the option is merely supererogatory. If an option is morally and overall best relative to all permissible weightings of moral and non-moral

¹⁷ In practice, *not* going on a bike ride typically also involves a tiny chance of causing harm to innocents, but I take it that it would be OK for me to go even if for some reason I could be sure that I wouldn't harm anyone by staying at home.

¹⁸ Alternatively, we might aggregate the moral and non-moral values only on the level of options, defining the permissible overall rankings of options as a weighted average of their expected moral value and their expected non-moral value. Mathematically, these two approaches are equivalent. Of course, further, more complicated approaches are also possible.

values, then it is obligatory.¹⁹

In practice, the division between the two kinds of values may be a little blurry, because the agent's personal goals and intentions often matter for the purely moral evaluation. Torturing and killing animals for fun is widely regarded as worse than torturing and killing animals for profit. (In many countries, the former is illegal and the latter subsidised.) A moral value function can certainly be sensitive to such intentional aspects of profiles. However, one might also try to account for the relevant judgements in terms of the balancing of moral and non-moral values. That is, one might assign moral disvalue to torturing and killing animals (ignoring motivation), and count profit, but not fun, as a good enough nonmoral reason to outweigh the moral disvalue. The two accounts are not equivalent. For example, they give different verdicts about the case of a person who tortures and kills animals, sells their meat for profit, but is motivated largely by the joy of killing and torturing.

8 GUIDANCE

I have assumed that moral theories need to evaluate the options in decision problems with imperfect information: they should answer questions about what an agent ought to do given such-and-such options and such-and-such information about the world. One motivation for this assumption was that moral theories should be “action guiding” in the sense that endorsing a moral theory should manifest itself in judgements about the choices people should make in real-life decision problems, where not all information is given. I have suggested that the framework of decision theory provides an attractive method to address this need, by deriving subjective moral verdicts from objective verdicts. The decision-theoretic approach, I argued, is in no way reserved for consequentialist theories. On the contrary, in some respects it is easier to defend on non-consequentialist grounds.

My proposal concerns the internal structure of moral theories. The proposal is that a theory's verdicts about decision problems with incomplete information should be related in a certain systematic way to its verdicts about decision problems with complete information.²⁰ It is not a proposal about the cognitive mechanisms involved in endorsing

¹⁹ Note that even if the moral value function is unique, we can now get a range of morally permissible options in many decision problems. So our solution to the problem of supererogation further helps with the problem of optionality (as one might have expected).

²⁰ This explains why I have paid no attention to moral uncertainty. What should you do, say, in a trolley problem if you are unsure whether it would be right or wrong to flip the switch? From the internal perspective of moral theories, the answer is usually clear. If a theory says that it would be wrong to flip the switch, it will presumably say that it would be wrong to do so no matter if the agent knows that it is. Merely believing that something is right or wrong doesn't make it so. (I don't rule out that a moral theory might have special rules for cases of moral uncertainty, although that won't help

a moral theory or in applying it to a given decision situation. In particular, my proposal is not that the cognitive system of morally ideal agents must somehow store a moral value function and then compute the expected value of her options on the basis of that function and the subjective probabilities. How to make cognitive systems of various kinds systematically follow the advice of decision theory is an interesting question, but it is not a question for ethics.

On the other hand, it is an open question whether it is physically and biologically possible for agents like us to always and reliably choose options that (say) maximize expected value, relative to some non-trivial value function. In real-life decision problems, the relevant space of profiles and options is typically very large, which makes it infeasible to directly store the answers for all decision problems, but also intractable to compute the answers within the time and energy constraints imposed by the decision situation (see e.g. [Gershman and Daw 2012]). Thus it is not implausible that decision-theoretic moral theories provide a standard of (subjective) rightness that agents like us cannot systematically live up to. One might argue that if the best feasible approximation to a general decision-theoretic choice procedure would deviate from the decision-theoretic ideal in certain situations, then what should (subjectively) be done in those situations is what the best approximation would do. I don't think this is correct, but even if it were, the resulting notion of subjective moral status would still be approximately decision-theoretic, so it would still vindicate decision-theoretic considerations in moral theories.

Some authors interpret the guidance requirement in a very different way. On that interpretation, a theory is action-guiding only if it provides a usable recipe for conscious deliberation that agents are supposed to follow when facing a choice. Decision theory is sometimes thought to put forward such a recipe, along something like the following lines: "First, create a mental list of all available options. For each option, consider all possible profiles that may come about by choosing the option. Multiply the moral value of each profile by its probability conditional on the option, and add up the products. The result is the expected value of the option. Then choose the option with greatest expected value." This is an obviously impractical and ridiculous way to address realistic decision situations, if only because the space of available options and profiles is generally huge.

Properly understood, decision theory is not about how agents should go about making choices. It does not prescribe a particular decision procedure, nor does it say that people should see expected utility maximization as a goal. Decision theory merely says *which* options an agent should choose, given their goals and information, not *how* they should choose these options.²¹ To be sure, we often face a choice not only about how to act,

agents who are uncertain whether the theory is true.)

21 This is implied by the orthodox approach to decision theory on which decision theoretic norms are ultimately norms on preferences, as well as by the functionalist accounts mentioned in note 5 above. It is also frequently emphasized by friends of decision theory; see, among many others, [Pettit 1991:

but also about how to decide how to act: say, whether to trust our gut, to follow some heuristics, or to draw up a decision matrix. Decision theory can be applied to these choices. It will recommend whatever method has greatest expected utility. Drawing up a complete decision matrix will very rarely come out as the best procedure (even if it is an available option, as it usually isn't), and so decision theory will generally advise against it.

In the same way, I take it that moral theories deal in the first place with choices between acts – say, whether or not one should flip the switch in a given trolley problem with imperfect information. The answer will leave open what mechanisms or procedures the agent is supposed to use in order to make her choice. Decision-theoretic moral theories (of the kind I advocate) therefore do not imply that agents should be actively “guided” by the principle to maximize expected moral value, in the sense of [Smith 2012]. As above, we may also consider choices between decision procedures, but these don't get a special treatment. If an agent faces a choice between deciding by different decision methods, decision-theoretic moral theories will rank the methods by their expected moral value. The best method will depend on the situation: sometimes it may be to follow one's gut, other times to apply some heuristic, and occasionally it may involve setting up a simplified decision matrix.²²

So decision theory does not prescribe any general procedure or recipe for making decisions. It fails the procedural guidance requirement – not because it prescribes an implausible procedure, but because it prescribes no general procedure at all. However, I don't think this is a flaw. Sensible moral theories shouldn't prescribe a general procedure for deciding between acts. For one thing, can't one always imagine situations in which using a given procedure would have horrible consequences? Even setting aside scenarios like these, the best procedure plausibly depends on contingent facts about the decision situation and the agent's cognitive system. If an agent happens to know that she has infallible moral intuitions which always favour acts with greatest expected value, she ought to follow those intuitions. If an agent is very good at maths and at introspecting her degrees of belief, computing expected values may be a useful procedure. If a decision has to be made within milliseconds, different procedures are advisable than if the agent has months to prepare. No single recipe could fit all decision-makers under all circumstances.²³

167–169], [Jackson 1991: 468–471], [Joyce 1999: 80].

²² The decision procedure that maximizes expected moral value can lead to a choice that doesn't maximize expected moral value. In that case, the agent did one thing right and another thing wrong: she chose the subjectively right decision procedure, but the subjectively wrong act.

²³ Even a very complicated disjunctive recipe, such as the one outlined in [Smith 2010], is problematic. For example, applying it may often take too long. Smith's recipe is more plausible not as a single disjunctive procedure, but as a disjunction of simpler procedures, applicable to different decision situations. In effect, decision-theoretic moral theories also determine such a disjunctive assignment, insofar as that they also recommend different procedures for different situations. Moreover, they

You might object that if an agent doesn't know, and can't compute, the expected values, it is of little help to say that she ought to use whatever decision procedure maximizes expected value. But recall that when I say that the agent should use whatever procedure maximizes expected value, I do not mean that she ought to explicitly compute the expected moral value of the available procedures. As before, decision theory merely prescribes a choice, not a procedure to arrive at that choice. (It just happens that the choice here is a choice between procedures.) You may still press your point: how in practice is the agent supposed to figure out what procedure to use? The question makes sense as a question about cognitive implementation: how might the agent's cognitive system end up making the decision-theoretically optimal choice? As I said, this is a good question, on which artificial intelligence and cognitive science have made considerable progress. But it is not a question for ethics. Alternatively, the question could also be meant as a question in applied ethics: what would be the morally right way for the agent to figure out a decision procedure? As such, the question presupposes that the agent could in principle use a variety of ways – that she faces a “third-order” choice between different ways to arrive at a decision procedure to choose between acts. Unsurprisingly, decision theory says that she ought to choose whatever second-order procedure maximizes expected moral value. And so on, for fourth-order choices and fifth-order choices. However, the hierarchy can't go on forever: at some point, the decision process must get started. The conscious mind is not an unmoved mover. At the highest level, there is no longer a moral question about how the agent is supposed to determine the right options.

REFERENCES

- Maurice Allais [1953]: “Le comportement de l'homme rationnel devant le risque: critique des postulats et axiomes de l'école américaine”. *Econometrica: Journal of the Econometric Society*: 503–546
- Allen R. Bernstein and Frank Wattenberg [1969]: “Non-Standard Measure Theory”. In W. Luxemburg (Ed.) *Applications of Model Theory of Algebra, Analysis, and Probability*, Holt, Reinhart and Winston
- Simon Blackburn [1998]: *Ruling Passions: A Theory of Practical Reasoning*. Oxford: Oxford University Press
- Ethan Bolker [1966]: “Functions Resembling Quotients of Measures”. *Transactions of the American Mathematical Society*, 124: 293–312

answer a question that becomes pressing in any account along Smith's lines: why exactly is procedure X adequate under conditions Y ? Is that a brute moral fact? In the decision-theoretic framework, we can answer that X is the right procedure under conditions Y because it then maximizes expected moral value.

- John Broome [1991]: *Weighing Goods*. Oxford: Blackwell
- Ruth Chang [2002]: “The Possibility of Parity”. *Ethics*, 112(4): 659–688
- Mark Colyvan, Damian Cox and Katie Steele [2010]: “Modelling the Moral Dimension of Decisions”. *Nous*, 44(3): 503–529
- James Dreier [1996a]: “Accepting Agent-Centered Norms: a Problem for Non-cognitivists and a Suggestion for Solving It”. *Australasian Journal of Philosophy*, 74: 409–421
- [1996b]: “Rational Preference: Decision Theory as a Theory of Practical Rationality”. *Theory and Decision*, 40: 249–276
- Andy Egan [2007]: “Some Counterexamples to Causal Decision Theory”. *Philosophical Review*, 116: 93–114
- Daniel Ellsberg [1961]: “Risk, Ambiguity and the Savage Axioms”. *Quarterly Journal of Economics*, 75: 643–669
- Heidi Li Feldman [1995]: “Science and uncertainty in mass exposure litigation”. *Texas Law Review*, 74: 1–48
- Samuel Gershman and Nathaniel D. Daw [2012]: “Perception, action and utility: the tangled skein”. In M. Rabinowich, K. Friston and P. Varona (Eds.) *Principles of Brain Dynamics: Global State Interactions*, Cambridge (MA): MIT Press, 293–312
- Joshua Gert [2004]: “Value and Parity”. *Ethics*, 114: 492–520
- Allan Gibbard and William Harper [1978]: “Counterfactuals and Two Kinds of Expected Utility”. In C.A. Hooker, J.J. Leach and E.F. McClennen (Eds.) *Foundations and Applications of Decision Theory*, Dordrecht: D. Reidel, 125–162
- William Harper [1984]: “Ratifiability and Causal Decision Theory: Comments on Eells and Seidenfeld”. *PSA*, 2: 213–228
- Frank Jackson [1991]: “Decision-Theoretic Consequentialism and the Nearest and Dearest Objection”. *Ethics*, 101(3): 461–482
- Frank Jackson and Michael Smith [2006]: “Absolutist Moral Theories and Uncertainty”. *The Journal of Philosophy*, 103: 267–283
- Richard Jeffrey [1965]: *The Logic of Decision*. New York: McGraw-Hill
- James Joyce [1999]: *The Foundations of Causal Decision Theory*. Cambridge: Cambridge University Press

- Seth Lazar [Unpublished]: “In Dubious Battle: Uncertainty and the Ethics of Killing”
- David Lewis [1974]: “Radical Interpretation”. *Synthese*, 23: 331–344
- [1979]: “Attitudes *De Dicto* and *De Se*”. *The Philosophical Review*, 88: 513–543
- [1981]: “Causal Decision Theory”. *Australasian Journal of Philosophy*, 59: 5–30
- [1989]: “Dispositional Theories of Value”. *Proceedings of the Aristotelian Society*, 63: 113–137. Reprinted in Lewis’s *Papers in Ethics and Social Philosophy*, Cambridge: Cambridge University Press, 2000
- Christopher J.G. Meacham and Jonathan Weisberg [2011]: “Representation theorems and the foundations of decision theory”. *Australasian Journal of Philosophy*, 89(4): 641–663
- Philip Pettit [1991]: “Decision Theory and Folk Psychology”. In M. Bacharach and S. Hurely (Eds.) *Foundations of Decision Theory: Issues and Advances*, Cambridge (MA): Blackwell, 147–175
- Douglas W. Portmore [2008]: “Are Moral Reasons Morally Overriding?” *Ethical Theory and Moral Practice*, 11(4): 369–388
- Wlodek Rabinowicz [2008]: “Value Relations”. *Theoria*, 74(1): 18–49
- Frank Ramsey [1931]: “Truth and Probability (1926)”. In R.B. Braithwaite (Ed.) *Foundations of Mathematics and other Essays*, London: Routledge & P. Kegan, 156–198
- Leonard Savage [1954]: *The Foundations of Statistics*. New York. Wiley
- Brian Skyrms [1984]: *Pragmatics and Empiricism*. Yale: Yale University Press
- Holly Smith [2010]: “Subjective Rightness”. *Social Philosophy and Policy*, 27(02): 64–110
- [2012]: “Using Moral Principles to Guide Decisions”. *Philosophical Issues*, 22(1): 369–386
- Jordan Howard Sobel [1983]: “Expected utilities and rational actions and choices”. *Theoria*, 49: 159–183. Reprinted with revisions in [Sobel 1994: 197–217]
- [1994]: *Taking Chances*. Cambridge: Cambridge University Press
- Robert Stalnaker [1981]: “Letter to David Lewis”. In W. Harper, R. Stalnaker and G. Pearce (Eds.) *Ifs: Conditionals, Belief, Decision, Chance, and Time*, Dordrecht: Reidel, 151–152

— [1984]: *Inquiry*. Cambridge (Mass.): MIT Press

Amos Tversky [1975]: “A critique of expected utility theory: Descriptive and normative considerations”. *Erkenntnis*, 9(2): 163–173

John von Neumann and Oskar Morgenstern [1944]: *Theory of games and economic behavior*. Princeton: Princeton University Press

Paul Weirich [1983]: “A decision maker’s options”. *Philosophical Studies*, 44(2): 175–186